
DocuVault

Knowledge Management AI Platform

Client Reference Guide

Version	DocuVault v2.0
Date	March 11, 2026
Prepared by	Renata AI
Classification	Confidential – Client Use Only

Contents

1. System Overview

2. Key Features

- 2.1 Document Management
- 2.2 AI Knowledge Assistant
- 2.3 Users, Roles & Access Control
- 2.4 Collaboration & Sharing
- 2.5 Search & Organisation
- 2.6 Audit Trail & Notifications

3. Accessing System Features

4. Architecture Overview

- 4.1 How the System is Structured
- 4.2 How Documents are Indexed
- 4.3 How the AI Answers Questions

5. Technical Reference

6. Add-On Features & Integrations

- 6.1 Annual Maintenance Contract (AMC)
- 6.2 Email System Integration
- 6.3 LAN Drive Sync (Google Drive-Style)
- 6.4 Auto-Indexing on Upload
- 6.5 Desktop Agent — Web to Desktop
- 6.6 Offline Operation — No Internet Required

1. System Overview

DocuVault is Renata AI's Knowledge Management Platform. It gives your organisation one place to store, organise, and query all your documents — using plain English, not search keywords.

At its core is an AI assistant that reads your documents and answers questions instantly. No more hunting through folders or opening file after file.

What you can do with DocuVault

- Upload and organise your documents in one secure place.
- Control who can see each document — from fully public to private or role-restricted.
- Ask the AI assistant questions in plain English and get answers drawn from your own documents.
- Collaborate with your team through comments, shared links, and notifications.
- Track every action with a complete audit trail.
- Manage users and roles with fine-grained permission levels.

Item	Details
Platform	DocuVault v2.0 — web-based (browser) + optional desktop agent
Supported Files	PDF, Word, text files, images, and more
Max File Size	100 MB per upload
AI Knowledge Modes	Hybrid (default) · Strict (documents only) · Indicated
Access Control	Public · Private · Role-Based · Custom (per user)
Web Access	Any modern web browser on desktop, tablet, or mobile
Desktop Access	Optional background desktop agent — auto-connects files, runs in system tray
Internet Required	No — fully offline after one-time model download. No API key, no cloud dependency.

2. Key Features

2.1 Document Management

Everything you need to manage documents throughout their full lifecycle.

Feature	What it does
Upload	Add documents from your computer. Supported types include PDF, Word, images, and more.
Access Control	Choose who can see each document: Public, Private, Role-Based, or shared with specific people.
Version History	Every time a document is edited, the previous version is saved. You can view or restore any past version.
Document Locking	Lock a document while editing so no one else can make changes at the same time.
Categories & Tags	Organise documents into categories (with sub-categories) and tag them for flexible grouping.
Metadata	Track view count, download count, file size, and upload date automatically.
Soft Delete	Deleted documents can be recovered by an administrator before permanent removal.

2.2 AI Knowledge Assistant

The AI assistant lets you ask questions about your documents in plain English. It reads through all indexed documents, finds the most relevant parts, and writes a clear answer.

Feature	What it does
Ask in Plain English	Type a question naturally — the AI understands context, not just keywords.
Answers from Your Docs	The AI draws from documents stored in DocuVault, so answers are grounded in your data.
Follow-up Questions	Ask follow-up questions in the same conversation. The AI remembers context.
Three Answer Modes	Hybrid (default): Uses your documents first, fills gaps with general knowledge. Strict: Only answers from your documents — refuses if the answer is not in any document. Indicated: Labels which parts come from your documents and which from general knowledge.
Permission-Aware	The AI only uses documents the logged-in user is allowed to access.
Source Transparency	Each answer shows which documents and pages were used, so you can verify the information.

2.3 Users, Roles & Access Control

User Type	What they can do
Guest	View documents that are set to Public only. Cannot upload or edit.
Regular User	Upload and manage their own documents. Access documents based on their role level. Comment and collaborate.
Admin	Full access to all documents, all users, and all system settings.

Administrators can create custom roles (e.g. "Team Lead", "Department Head") and assign a numeric level to each. Documents set to Role-Based access are visible only to users whose role level meets or exceeds the requirement.

2.4 Collaboration & Sharing

Feature	What it does
Comments	Add comments to any document. Replies are threaded for easy reading.
Shared Links	Generate a link to share a document. Optionally add a password, set an expiry date, or limit how many times it can be opened.
Direct Sharing	Share a document directly with specific registered users.
Notifications	The system automatically alerts you when a document is shared with you, updated, or commented on.
Favourites	Bookmark documents you use often for quick access from your sidebar.

2.5 Search & Organisation

Feature	What it does
Search Bar	Search by title, description, or content. Results rank the most relevant documents first.
Filters	Narrow results by owner, date, access level, category, tags, or file type.
Sorting	Sort by title, upload date, last updated, view count, or file size.
Categories	Organise documents into a tree of categories. Each category can have a colour and icon.

Tags	Add one or more tags to a document. Search and filter by tag across all categories.
------	---

2.6 Audit Trail & Notifications

Every action in DocuVault is automatically recorded. The activity log shows who did what, and when — making it easy to track changes for compliance and accountability.

Recorded Action	When it appears in the log
Upload / Create	A new document is added to the system.
View / Download	A user opens or downloads a document.
Edit / Update	A document or its details are changed.
Delete	A document is removed (soft-deleted).
Share	A document is shared via link or direct user share.
Comment	A comment is added to a document.
Permission Changed	A document's access level or user role is updated.

3. Accessing System Features

The table below shows how to navigate to each feature. All paths are relative to your system's base URL (e.g. <https://yourdomain.com>).

URL Path	How to access
/register/	Open the system in your browser and click Register. Enter your name, email, and password.
/login/	Enter your username and password and click Sign In.
/dashboard/	Your home screen after login. Shows recent documents, notifications, and quick links.
/documents/	The full document library. Use the search bar and filter panel to narrow results.
/documents/create/	Click + New Document. Select a file, fill in the details, set access level, and click Save.
<a href="/documents/<id>/">/documents/<id>/	Click any document title to open it — preview, download, comment, version history, and share.
<a href="/documents/<id>/edit/">/documents/<id>/edit/	Open a document then click Edit. Save changes — the system automatically creates a new version.
<a href="/documents/<id>/index/">/documents/<id>/index/	Open a document then click Index for AI. The system reads and indexes the content for the AI assistant.
/documents/bulk-index/	Admins only. From the document list, go to Actions → Bulk Index to index all unprocessed documents.
/chatbot/	Click AI Assistant in the sidebar. Type your question and press Enter.
/search/	Use the search bar in the top navigation. Add filters for date, owner, category, tags, and access level.
/categories/	Navigate to Organise → Categories to create, nest, and colour-code categories.
/favorites/	Click the star icon on any document. Access bookmarks via Favourites in the sidebar.
/notifications/	Click the bell icon in the top bar to see all notifications.
/activity/	Go to Account → Activity Log for a complete history of all actions.
/admin/users/	Admins only. Go to Admin → Users to view all accounts and update role assignments.
/admin/roles/	Admins only. Go to Admin → Roles to create, edit, or remove custom roles.
/profile/edit/	Click your avatar (top right) → Edit Profile to update your details.

4. Architecture Overview

This section gives a plain-English view of how DocuVault is structured — how it stores your documents, how the AI learns from them, and how it answers your questions. Every component described here runs on your own hardware with no internet connection required.

4.1 How the System is Structured

DocuVault has three main layers that work together:

Layer	What it is	What it does for you
Your Browser	The web interface you open on any device.	Where you upload, search, chat with AI, and manage settings.
DocuVault Platform	The server application in the background.	Handles login, stores documents securely, controls access, and connects you to AI.
AI & Storage	The AI engine and databases.	Reads and indexes documents, answers questions, and stores all data.

System Flow




Figure 1: All components run on your local server — no internet required during operation.

4.2 How Documents are Indexed

Before the AI can answer questions about a document, it must be "indexed" — this means the AI reads it and saves a summary in a way it can search very quickly.




Figure 2: The five steps from uploading a document to it being ready for AI questions.

Steps in plain English

Step 1 — Upload: You select a file and upload it. DocuVault saves it securely.

Step 2 — Validate: The system checks the file size (must be under 100 MB) and file type.

Step 3 — AI Reads It: The platform reads the document — extracting all text, tables, and if needed, running text recognition on scanned pages.

Step 4 — Save to Index: The content is broken into sections and saved in the AI search index. This lets the AI find relevant parts instantly.

Step 5 — Ready: The document is now available to the AI assistant for any permitted user.

How to index a document

Open any document, then click Index for AI. The status will update to "Indexed" once complete. Admins can also index all documents at once using the Bulk Index option.

4.3 How the AI Answers Questions

When you type a question in the AI assistant, DocuVault goes through these steps to find and write the answer:




Figure 3: How a question becomes an answer.

Steps in plain English

Step 1 — You type a question: Ask anything in plain English, e.g. "What are the leave policy rules?"

Step 2 — Platform searches your documents: DocuVault searches all indexed documents to find the sections most relevant to your question.

Step 3 — AI reads the relevant parts: The AI reads the top matching sections. It only uses documents you have permission to access.

Step 4 — AI writes the answer: Using the content it found, the AI writes a clear, direct answer.

Step 5 — You receive the answer: The answer appears in the chat with the source documents and page numbers used. You can ask follow-up questions.

The three answer modes

Mode	How it behaves	Best for
Hybrid	Uses your documents first. Fills gaps with general knowledge if needed.	General use — recommended for most teams.
Strict	Only answers from your documents. Says so clearly if information is not found.	Compliance, legal, or regulated content.
Indicated	Uses both sources but clearly labels which part comes from where.	Research, auditing, or when traceability matters.

5. Technical Reference

Technology Stack

Category	Technology
Web Framework	Django 5.x (Python 3.10+)
Database	SQLite (development) / PostgreSQL (production)
AI Language Model	Ollama (local inference) — llama3.1:8b or compatible model Downloaded once, runs fully offline. No API key required.
AI Search Engine	ChromaDB vector database — stored locally, no cloud dependency
Embedding Model	all-MiniLM-L6-v2 — runs on-device via sentence-transformers
Document Reading	pdfplumber, PyMuPDF, Camelot (tables), Tesseract (OCR for scanned pages)
AI Framework	LangChain + Ollama backend, HuggingFace (local), PyTorch
Max File Upload	100 MB per file
Search Method	Hybrid — 70% semantic similarity + 30% keyword matching
Internet Required	No — entire stack runs offline on local hardware after setup

AI Configuration

Setting	Value
LLM runtime	Ollama — local model server, starts automatically with DocuVault
Default model	llama3.1:8b (downloaded once, ~4.7 GB, then fully offline)
API key required	None — no Groq, OpenAI, or any cloud API key needed
Document chunk size	512 characters (256 in lightweight mode)
Overlap between chunks	100 characters
Results per query	8 chunks (6 in lightweight mode)
Conversation memory	Up to 8 turns per session
Max response length	512 tokens (~380 words)
Response consistency	Low variability — temperature 0.2 for reliable, repeatable answers

Internet during inference	Not required — model runs entirely on local CPU/GPU
---------------------------	---

User & Permission Model

Access Level	Who can see the document
Public	Everyone, including guests who are not logged in.
Private	The document owner and admins only.
Role-Based	Users whose role level is equal to or higher than the required level set by the owner.
Custom	Only specific users selected by the document owner, plus admins.

Key Database Records

Record Type	What it stores
Document	The file, title, access level, category, tags, and owner.
Document Version	A snapshot of the document every time it is edited.
Chat Session	A conversation thread between a user and the AI assistant.
Chat Message	Each individual question and AI answer, with source references.
Activity Log	An immutable record of every action in the system.
Notification	Alerts sent to users for shares, comments, and updates.
Shared Link	A temporary link with optional password, expiry, and access count.

6. Add-On Features & Integrations

DocuVault is built to grow with your organisation. The features described in this section are available as add-ons or integrations that can be configured and enabled on top of the standard platform — either during initial deployment or at any later stage.

6.1 Annual Maintenance Contract (AMC)

An **Annual Maintenance Contract (AMC)** can be provided to keep DocuVault updated, secure, and running smoothly over time. Through AMC, the system can receive improvements such as newer AI models, performance optimisations, security updates, and feature enhancements without affecting your existing documents.

What the AMC covers

- **AI model upgrades** – The system can be updated to newer AI models that provide better understanding, improved responses, and more accurate answers.
- **Performance improvements** – Enhancements to document search, response speed, and overall system efficiency.
- **Security updates and monitoring** – Regular updates to maintain system security and stability.
- **UI improvements and new features** – Periodic updates to improve usability and add useful capabilities.

Important clarification

DocuVault does **not train the AI model on your documents**. Instead, it uses an approach called **Retrieval-Augmented Generation (RAG)**.

In simple terms, three components work together in the system:

AI Model (running through Ollama)

This is the AI engine that understands questions and generates answers in natural language.

Embedding Model

This model converts document text into a searchable format so the system can quickly find the most relevant parts of documents.

RAG Process

When a user asks a question, the system first searches the documents using the embeddings, retrieves the most relevant sections, and then the AI model reads those sections and generates a clear answer.

Because of this design:

- **AI model upgrades do not require retraining on your documents.**
- **Your existing indexed documents will continue to work normally after an upgrade.**

In simple terms, the **documents stay the same, but the AI engine that interprets them becomes more advanced.**

When re-indexing may be needed

Re-indexing is only required if the **embedding model used for document search changes**.

If that happens, the system simply **regenerates the document embeddings automatically**. The original documents do not need to be uploaded again, and this process can be handled as part of the AMC.

6.2 Email System Integration

DocuVault can integrate with your organisation's email system. Emails and their attachments are ingested and indexed so they become searchable through the AI assistant — just like any other document.

Supported email systems

- Microsoft Outlook / Exchange
- Microsoft 365
- Gmail
- IMAP servers
- On-premise enterprise mail servers

What gets indexed

Email Content	What is indexed
Email body	The full text of the email message.
Attachments	PDF, Word, Excel, and PowerPoint files attached to emails.
Metadata	Sender, receiver, subject line, and timestamp.

Security & access control

Role-based access ensures that emails are only retrievable by users who are authorised to see them.

- HR emails → accessible to HR users only
- Finance emails → accessible to Finance users only

Once integrated, users can ask the AI assistant questions like: "Show discussions about vendor contracts with ABC company." The system retrieves matching content from documents, attachments, and email conversations together.

Email integration flow




Figure 4: How emails flow from your mail server into DocuVault and become searchable via the AI assistant.

6.3 LAN Drive Sync (Google Drive-Style)

Instead of manually uploading files through the web interface, users can save documents into a synchronised folder on their computer. A background sync agent watches the folder and automatically sends new or updated files to DocuVault.

How it works




Figure 5: LAN sync pipeline — from a user's folder to DocuVault version control and AI indexing.

Features

- Automatic document synchronisation — no manual uploads required.
- Version tracking — every save creates a new version with timestamp and user.
- Rollback — restore any previous version from DocuVault.
- Instant AI indexing — new documents are indexed automatically as they sync.
- Enterprise-ready — supports multiple PCs on the same LAN.

Implementation options

Option	Description
Desktop sync agent (recommended)	A lightweight background app installed on each user's PC.
Network shared drive monitoring	Monitors a shared network folder without any client install.
Drive-style desktop client	A full Google Drive-like interface with sync status indicators.

6.4 Auto-Indexing on Upload

By default, documents must be manually triggered for AI indexing after upload. The system can be configured to index documents automatically the moment they are uploaded, so they are immediately available to the AI assistant without any extra step.

How indexing works (not AI training)

A common question from clients is whether uploading documents "trains" the AI. The answer is no. DocuVault uses indexing, not training. The distinction is important:

	AI Training	DocuVault Indexing
What it is	Modifying the AI model's internal weights.	Reading and storing document content in a searchable format.
Time	Hours to days.	Seconds to minutes per document.
Effect	Permanently changes the AI model.	Adds document to the search index only.
Required?	Never required in DocuVault.	Required once per document (or on re-upload).

Indexing pipeline diagram




Figure 6: Left — Document processing pipeline (upload to vector database). Right — Query handling (user question to AI answer).

How to enable auto-indexing

Auto-indexing can be enabled by your system administrator in the platform configuration. Once enabled, every document uploaded will be automatically indexed in the background without any action required from the user.

6.5 Desktop Agent — Web to Desktop

By default, DocuVault runs entirely in a web browser — no installation needed. However, for teams who work heavily with local files, a lightweight Desktop Agent can be installed alongside the web platform. The agent runs quietly in the background, bridges your computer's file system directly to DocuVault, and makes accessing the platform feel like a native desktop app.

What the Desktop Agent does

Capability	Description
System tray icon	The agent sits in your Windows/macOS system tray. One click opens DocuVault instantly — no browser URL to remember.
Automatic file watching	Select one or more folders on your computer. Any file saved there is automatically detected and synced to DocuVault.
Background sync	Files are uploaded and indexed in the background without interrupting your work.

Auto-connect on startup	The agent starts with your computer and reconnects to the DocuVault server automatically — no manual login each day.
Local file shortcuts	Right-click any file in Windows Explorer / macOS Finder and choose "Send to DocuVault" directly.
Offline queue	If the server is temporarily unreachable, the agent queues pending files and syncs them automatically once the connection is restored.
Notification alerts	Desktop pop-up notifications for completed uploads, indexing status, shared documents, and AI query replies.
Session persistence	Stay logged in across sessions — the agent maintains your authentication token securely.

Web vs Desktop — side by side

	Web Browser (Standard)	Desktop Agent (Add-On)
Installation	None — open any browser	One-time lightweight install (~20 MB)
File access	Manual upload via browser	Automatic — watches selected folders
Startup	Open browser and navigate to URL	Starts with computer, always ready
Notifications	In-browser only	Native desktop pop-ups
Right-click upload	Not available	Right-click any file to send to DocuVault
Offline handling	Not available	Queues files and syncs when back online
Best for	Occasional users, mobile, tablet access	Power users who work with many local files daily

How the Desktop Agent connects




Figure 7: From saving a file on your computer to it being searchable in DocuVault — fully automatic.

Deployment

Supported OS	Windows 10/11, macOS 12+, Ubuntu 20.04+
Server connection	Connects to DocuVault over LAN or internet (HTTPS)
Authentication	Uses the same username / password as the web platform

Configuration	Admin sets allowed folders, sync rules, and notification preferences via a simple settings panel
Distribution	Installer provided by Renata AI — deployable via Group Policy (Windows) or MDM (macOS)

Who should use the Desktop Agent?

The Desktop Agent is ideal for users who regularly save documents locally (reports, drawings, scanned files, spreadsheets) and want them available in DocuVault without any manual effort. The web platform remains fully functional alongside it — both modes work together.

6.6 Offline Operation — No Internet Required

DocuVault is designed to run entirely on your local network or on a single machine with no internet connection required at any point during normal operation. The AI model, the search engine, the database, and the file storage all run on your own hardware — your data never leaves your premises.

What runs locally

Component	How it runs offline
AI Language Model (LLM)	Ollama runs llama3.1:8b (or any compatible model) directly on your server CPU or GPU. Model is downloaded once during setup — no API calls made during use.
Embedding Model	all-MiniLM-L6-v2 runs on-device via sentence-transformers and PyTorch. No external API needed.
Vector Database	ChromaDB stores all document embeddings as local files on the server disk.
Document Storage	All uploaded files are stored in the local file system under /media/.
Web Application	Django serves the web interface on your local network. Users access it via LAN IP or a local domain — no internet required.
Database	SQLite or PostgreSQL runs locally. All user data, roles, and activity logs stay on-site.
Desktop Agent	Connects to the DocuVault server over LAN (HTTPS). Never calls any external service.

Setup — one-time steps (internet needed once only)

Step 1 — Install DocuVault: Install the platform on your server. This requires internet to download the software package once.

Step 2 — Download the AI model: Run: `ollama pull llama3.1:8b` (or the model specified by Renata AI). This downloads the model file (~4.7 GB) to the server. Done once — never again.

Step 3 — Download embedding model: The all-MiniLM-L6-v2 embedding model (~90 MB) is downloaded automatically on first run and cached locally.

Step 4 — Go fully offline: Once both models are cached, disconnect from the internet. DocuVault continues to operate with full AI functionality indefinitely.

Offline capability summary

Feature	Works Offline?	Notes
Document upload & management	Yes	Files saved locally on server.
AI assistant (Q&A)	Yes	Ollama runs the model locally — no API call.
Document indexing	Yes	Embeddings generated on-device.
User login & access control	Yes	Authentication handled by local Django server.
Search & filters	Yes	ChromaDB vector search runs locally.
Email integration	Yes	Connects to on-premise mail server on LAN.
Desktop Agent sync	Yes	Communicates with DocuVault server over LAN.
LLM model updates	Requires internet once	New model pulled via ollama pull, then offline again.
Software updates	Requires internet once	Update package downloaded, then offline again.

Hardware recommendations for offline deployment

Component	Minimum	Recommended
CPU	4 cores (model runs on CPU)	8+ cores for faster AI responses
RAM	16 GB	32 GB for smooth multi-user operation
Storage	50 GB free	200 GB+ for large document libraries
GPU	Not required	NVIDIA GPU (4 GB+ VRAM) for significantly faster AI responses
OS	Ubuntu 20.04+ / Windows Server	Ubuntu 22.04 LTS (recommended)

Network	LAN only (no internet needed)	1 Gbps LAN for fast file sync
---------	-------------------------------	-------------------------------

Your data stays on your premises — always

Because DocuVault runs entirely on your own hardware with no cloud dependency, your documents, queries, and AI responses are never transmitted to any external server. This makes DocuVault suitable for organisations with strict data-sovereignty, confidentiality, or air-gapped network requirements.

7. Client Q&A Summary

The following table summarises the most common questions raised by clients during technical discussions about DocuVault.

Client Question	Answer
Can you offer an AMC for LLM upgrades?	Yes. An AMC covers LLM upgrades, performance improvements, security patches, and UI enhancements.
Do you need to retrain the AI when upgrading the LLM?	No. DocuVault uses RAG — the LLM is not trained on documents. Upgrades do not affect indexed data.
When is re-indexing required?	Only if the embedding model itself is changed (e.g. switching providers). This is handled automatically under AMC.
Can DocuVault integrate with our email system (Outlook, Exchange)?	Yes. Emails and attachments are ingested and indexed. Users can query email content through the AI assistant.
Can it work like Google Drive with automatic sync?	Yes. A LAN Sync Connector can be implemented so documents sync automatically from user folders to DocuVault.
Does the AI train automatically when I upload a document?	No. The AI does not train on documents. Uploading triggers indexing, which is fast and does not change the AI model.
How is DocuVault different from BigQuery?	BigQuery is designed for structured data analytics using SQL. DocuVault is designed for semantic search and AI question answering over unstructured documents such as PDFs, emails, and reports.
Do users have to open a browser every time?	No. With the optional Desktop Agent, DocuVault runs in the system tray. Files sync automatically from watched folders, and the platform is accessible with a single click.
Can the Desktop Agent work if the server is offline?	Yes. The agent queues pending files locally and syncs them automatically once the server is reachable again.
Does the AI require an internet connection or an API key?	No. The AI runs fully offline using Ollama on your local server. No Groq, OpenAI, or any cloud API key is needed. The model is downloaded once during setup and then runs without internet.
Does data ever leave our network?	Never. All documents, queries, AI responses, and user data stay entirely on your own hardware. DocuVault has no cloud dependency during operation.
What happens if our internet goes down?	Nothing — DocuVault continues to work exactly as normal. All components (web app, AI model, database, file storage) run on your local server.
How much disk space does the AI model need?	The default model (llama3.1:8b) requires approximately 4.7 GB. This is downloaded once. Additional models can be added without affecting existing operation.